

ML Engineer Assessment — RAG System for AI Governance Documents

A Retrieval-Augmented Generation (RAG) system that answers questions over a corpus of AI governance documents. Retrieves relevant passages using hybrid search, fuses results with Reciprocal Rank Fusion, and generates grounded answers with source citations using Gemini 2.5 Flash.

How It Works

flowchart TD

Q[User Query]

D[Dense Retrieval
ChromaDB Similarity Search
Top-20]

S[Sparse Retrieval
BM25 Keyword Search
Top-20]

M[MMR Retrieval
ChromaDB MMR Search
Top-20]

R[RRF Fusion
Reciprocal Rank Fusion
k = 60]

F[MMR Diversification
Select Top-8 Diverse Chunks]

C[Context Assembly]

L[Gemini 2.5 Flash
Temperature = 0
Context-Only Prompt]

A[Grounded Answer
+
Source Citations]

Q --> D

Q --> S

Q --> M

D --> R

S --> R

M --> R

R --> F

F --> C

C --> L

L --> A

Ingestion pipeline:

flowchart LR

D[Documents]

C[RecursiveCharacterTextSplitter
Chunk Size = 1000
Overlap = 200]

E[all-MiniLM-L6-v2
Embeddings]

V[ChromaDB
Vector Store]

D --> C

C --> E

E --> V

Key Decisions

Hybrid retrieval (dense + sparse + MMR)

Dense retrieval captures semantic similarity but misses exact terminology (policy names, acronyms). BM25 fills that gap with keyword matching. The third MMR retriever adds a diverse set of semantically relevant candidates. All three are fused with RRF rather than picking one, which is more robust than any single retrieval strategy.

RRF over learned fusion

RRF (Reciprocal Rank Fusion) requires no training data and consistently outperforms simple score averaging in practice. The $k=60$ constant from the original paper dampens rank sensitivity — a document appearing at rank 1 in all three lists scores significantly higher than one appearing at rank 5.

Post-RRF MMR diversification

After fusion we have a single ranked list that can still contain semantically redundant chunks. A final MMR pass ($\lambda=0.5$) balances relevance against diversity, ensuring the 8 context chunks cover different facets of the answer rather than repeating the same passage.

ChromaDB over a managed vector store

Lightweight, local, and zero infrastructure overhead. Appropriate for assessment-scale workloads (11,625 chunks). Not a fit for production-scale or distributed ingestion.

all-MiniLM-L6-v2 as the embedding model

384-dimensional embeddings, fast indexing, strong baseline performance. BGE-M3 would give better multilingual and sparse support but is $\sim 15\times$ larger — the tradeoff isn't worth it for a single-language corpus.

Temperature=0 + context-only system prompt

Deterministic outputs and a hard constraint to only use retrieved context. If the corpus doesn't contain

evidence, the model responds with a fixed fallback rather than hallucinating.

Limitations

- **No reranker.** A cross-encoder reranker (e.g. bge-reranker-base) applied after RRF would improve answer quality but adds latency and model weight.
 - **BM25 is rebuilt per query.** The BM25 index is constructed in-memory from ChromaDB's stored chunks on each call. Acceptable at this scale; would need caching for production.
 - **ChromaDB is not production-grade.** Single-node, no distributed indexing, limited throughput.
 - **Chunk metadata is sparse.** The dataset provides plain text files without embedded page numbers or section headings, so source citations reference the file name only.
 - **No query rewriting.** Ambiguous or multi-hop questions are not decomposed before retrieval.
-

Setup

Prerequisites: Python 3.11+, a Google AI Studio API key (aistudio.google.com/apikey)

```
git clone https://github.com/arehh04/ml-assessment.git
cd ml-assessment

pip install -r requirements.txt

cp .env.example .env
# Edit .env and set GOOGLE_API_KEY
```

Download the dataset (requires a Kaggle account):

```
pip install kagglehub
python -c "
import kagglehub, shutil, pathlib
src = pathlib.Path(kagglehub.dataset_download('umerhaddii/ai-governance-
documents-data'))
dst = pathlib.Path('data/documents')
dst.mkdir(parents=True, exist_ok=True)
for f in (src / 'agora' / 'fulltext').iterdir():
    shutil.copy(f, dst)
print(f'Copied {len(list(dst.iterdir()))} documents')
"
```

Ingest documents:

```
python src/main.py --ingest
# Indexed 11625 chunks into ChromaDB at ./chroma_db
```

Running Queries

```
python src/main.py --query "Your question here"
python src/main.py --query "Your question here" --top-k 5    # use 5 context
chunks instead of 8
```

Running Tests

```
pytest tests/ --ignore=tests/test_integration.py    # 58 unit tests, no API key
needed
pytest tests/test_integration.py -m integration     # end-to-end, requires
dataset + API key
```

Test Queries

Query 1: Core topic

```
$ python src/main.py --query "What are the key principles of AI governance?"
```

Answer:

Key principles of AI governance include:

- * Adopting a common set of ethical AI principles, such as those based on the OECD AI Principles, to assist regulators and stakeholders and nurture responsible AI innovation.
- * Ensuring AI systems are robust, secure, and safe throughout their lifecycle, functioning appropriately and not posing unreasonable safety or security risks.
- * Establishing accountability for the proper functioning of AI systems and adherence to principles.
- * Implementing governance processes that catalog, review, and approve rights-impacting and safety-impacting uses of AI prior to deployment.
- * Cultivating a framework that fosters safe, secure, and trustworthy AI development, manages AI risks, respects human rights, civil rights, civil liberties, and privacy.

Sources:

- data\documents\1076.txt
- data\documents\1026.txt
- data\documents\1152.txt

- data\documents\1377.txt
- data\documents\764.txt
- data\documents\1396.txt
- data\documents\1356.txt
- data\documents\1385.txt

Query 2: Specific policy domain

```
$ python src/main.py --query "How should the Department of Defense govern AI systems?"
```

Answer:

The Department of Defense should govern AI systems by:

- * Improving coordination on AI governance with the defense industry sector.
- * Establishing and overseeing AI guidance on ethical requirements and protections,
and reducing or mitigating instances of unintended bias in AI algorithms.
- * Reviewing AI program funding to ensure DoD investment adheres to all established
AI policy.
- * Taking appropriate measures to ensure responsible development, deployment,
and use
of military AI capabilities throughout their lifecycle.
- * Prioritizing research to improve the security, robustness, and reliability of AI
systems and controls, and partnering with agencies, industry, civil society,
and academia.
- * Developing and implementing an AI training program for the covered workforce.

Sources:

- data\documents\403.txt
- data\documents\1380.txt
- data\documents\148.txt
- data\documents\1385.txt
- data\documents\399.txt
- data\documents\772.txt
- data\documents\27.txt

Query 3: Out-of-corpus question (hallucination mitigation)

```
$ python src/main.py --query "What does GDPR say about automated decision making?"
```

Answer:

I cannot find sufficient evidence in the provided documents to answer this question.

The corpus covers US federal AI policy and OECD principles — GDPR is not represented. The system correctly declines rather than fabricating an answer.

Future Works

Given more time:

1. **Cross-encoder reranker** (bge-reranker-base) applied after RRF for better precision
2. **Query rewriting** to decompose multi-hop questions before retrieval
3. **Evaluation pipeline** — Recall@K, MRR, and faithfulness scoring against ground truth
4. **Metadata enrichment** — extract document titles and section headings during ingestion for richer citations
5. **Persistent BM25 index** — serialise and cache to avoid rebuilding on every query